

GPU Project : Batch merge and merge path sort

Arthur **Zucker** , Clément **Apavou**



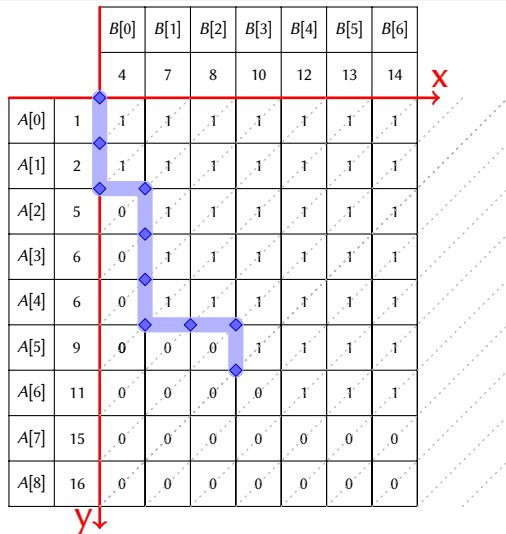
14 décembre 2020

Table of contents

- 1 Introduction
- 2 Merge path and sort
- 3 Batch merge
- 4 Merge sort applications


Introduction

Path merge : How it works



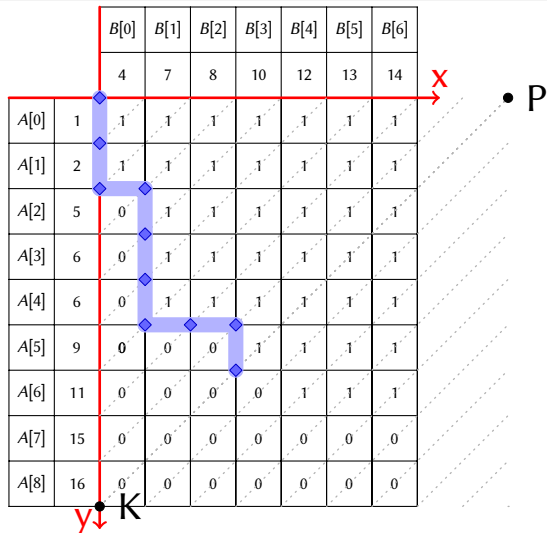
1 Compute P and K :

Path merge : How it works




1 Compute P and K :
→ $P=(9,0)$

Path merge : How it works



- 1 Compute P and K :
 $\rightarrow P=(9,0)$
 $\rightarrow K=(0,9)$

Path merge : How it works



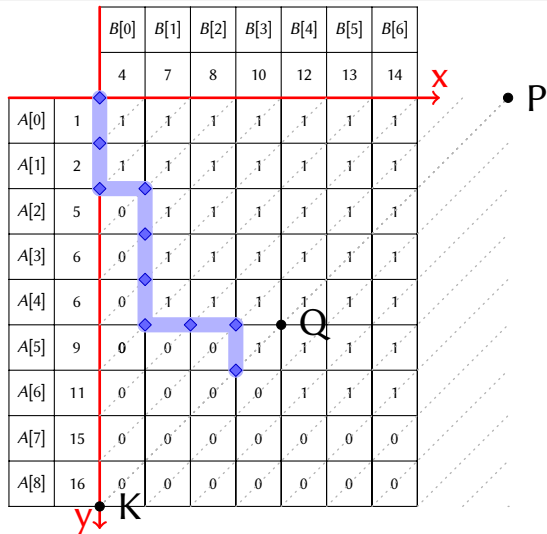
1 Compute P and K :

→ $P=(9,0)$

→ $K=(0,9)$

2 offset = 3

Path merge : How it works



1 Compute P and K :

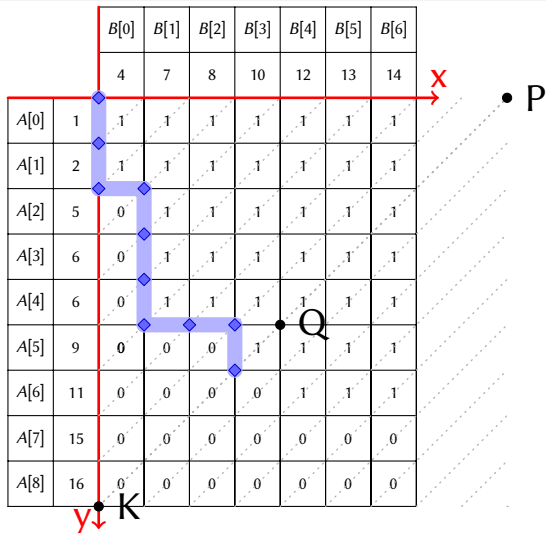
→ $P=(9,0)$

→ $K=(0,9)$

2 offset = 3

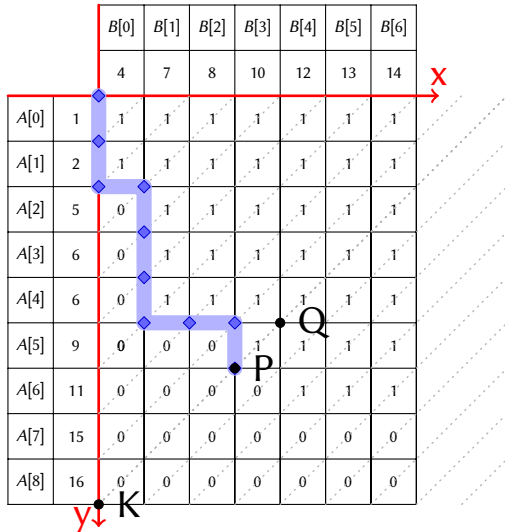
3 $Q = (5,4)$

Path merge : How it works



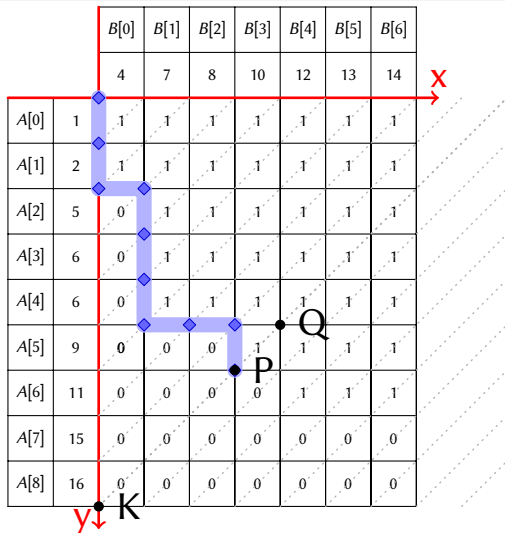
- 1 Compute P and K :
→ $P=(9,0)$
→ $K=(0,9)$
- 2 offset = 3
- 3 $Q = (5,4)$
- 4 $P = (Q_x-1, Q_y+1)=(4,5)$

Path merge : how it works



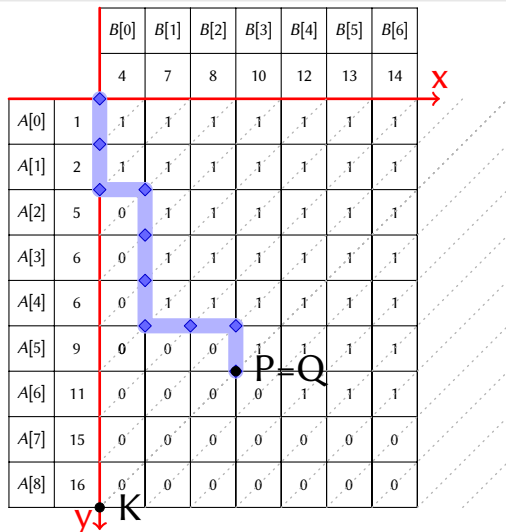
- 1 Compute P and K :
→ $P=(9,0)$
→ $K=(0,9)$
- 2 offset = 3
- 3 $Q = (5,4)$
- 4 $P = (Q_x-1, Q_y+1)=(4,5)$

Path merge : how it works



- 1 Compute P and K :
→ $P=(9,0)$
→ $K=(0,9)$
- 2 offset = 3
- 3 $Q = (5,4)$
- 4 $P = (Q_x-1, Q_y+1)=(4,5)$
- 5 $Q = P = (4,5)$

Path merged : thread 9



- 1 Compute P and K :
 $\rightarrow P=(9,0)$
 $\rightarrow K=(0,9)$
- 2 offset = 3
- 3 $Q = (5,4)$
- 4 $P = (Q_x-1, Q_y+1)=(4,5)$
- 5 $Q = P =(4,5)$

Path merged : thread 9, path

		B[0]	B[1]	B[2]	B[3]	B[4]	B[5]	B[6]
		4	7	8	10	12	13	14
A[0]	1	1	1	1	1	1	1	1
A[1]	2	1	1	1	1	1	1	1
A[2]	5	0	1	1	1	1	1	1
A[3]	6	0	1	1	1	1	1	1
A[4]	6	0	1	1	1	1	1	1
A[5]	9	0	0	0	1	1	1	1
A[6]	11	0	0	0	0	1	1	1
A[7]	15	0	0	0	0	0	0	0
A[8]	16	0	0	0	0	0	0	0

Path merged : thread 9, path

The diagram illustrates the process of removing the first row and column from an 8-bit grayscale image to extract bit-planes. It consists of two 9x8 grids, A and B.

- Grid B:** Contains the original bit values for each pixel. The columns are labeled $B[0]$ through $B[6]$. The rows are indexed 0 to 8. The values in the first row are 4, 7, 8, 10, 12, 13, 14.
- Grid A:** Shows the result after removing the first row and column. The first row and column of grid A contain zeros. The remaining elements are shifted up and to the left by one position relative to their original positions in grid B.

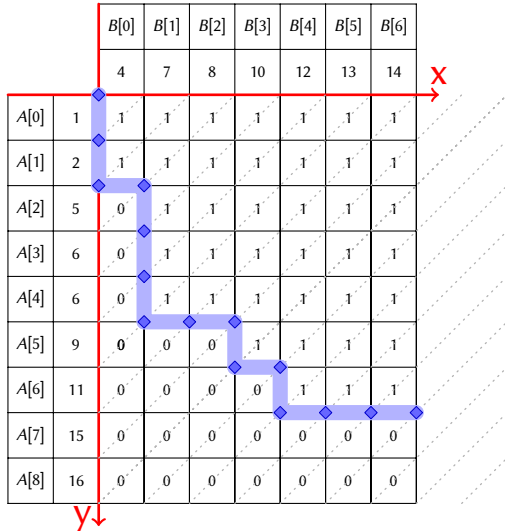
Blue diamonds mark the positions where the bit values have been shifted. Red arrows labeled **X** and **Y** indicate the removal of the first row and column, respectively.

Path merged : thread 9, path

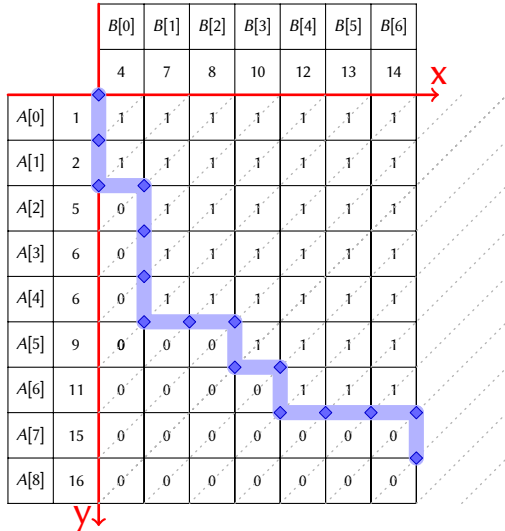
		$B[0]$	$B[1]$	$B[2]$	$B[3]$	$B[4]$	$B[5]$	$B[6]$
		4	7	8	10	12	13	14
$A[0]$	1	1	1	1	1	1	1	1
$A[1]$	2	1	1	1	1	1	1	1
$A[2]$	5	0	1	1	1	1	1	1
$A[3]$	6	0	1	1	1	1	1	1
$A[4]$	6	0	1	1	1	1	1	1
$A[5]$	9	0	0	0	1	1	1	1
$A[6]$	11	0	0	0	0	1	1	1
$A[7]$	15	0	0	0	0	0	0	0
$A[8]$	16	0	0	0	0	0	0	0



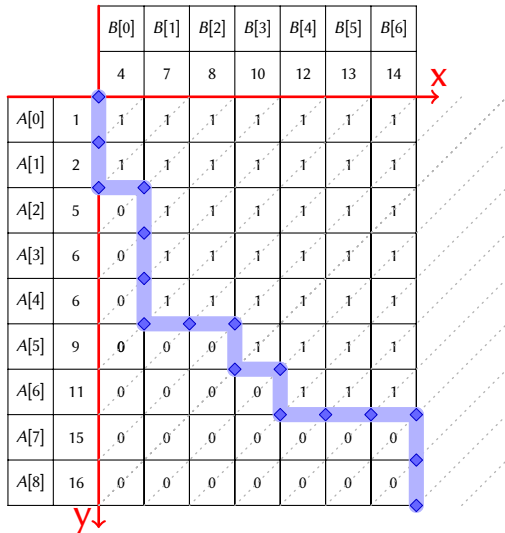
Path merged : thread 9, path



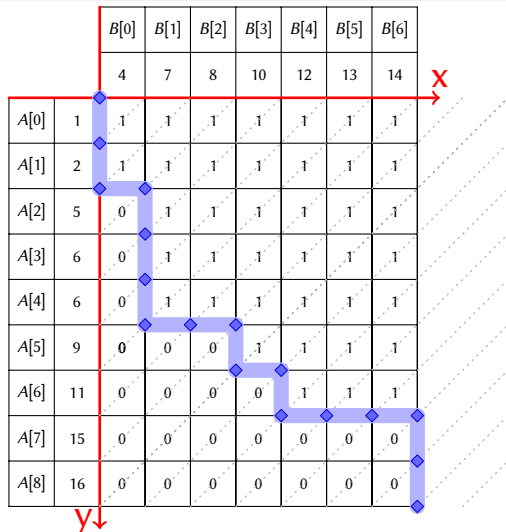
Path merged : thread 9, path



Path merged : thread 9, path



Path merged : thread 9, path



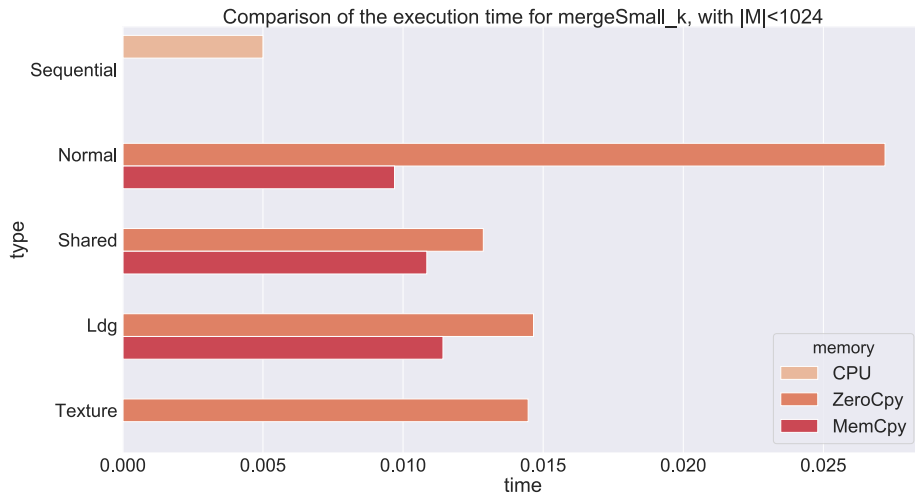
Merge path and sort

Question 1 : mergeSmall_k

Concept

- Two sorted arrays A and B such as : $|A| + |B| \leq 1024$
- mergeSmall_k merges A and B
- Using one block of threads

Question 1 : Different type of memory


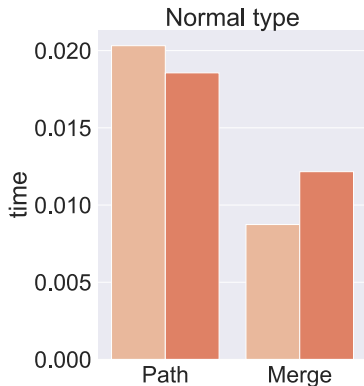


Question 2 : pathBig_k and mergeBig_k

Concept

- Two sorted arrays A and B such as : $|A| + |B| = d$
- pathBig_k finds the merge path
- mergeBig_k merges A and B

Question 2 : Different type of memory and methods



memory


- MemCpy
- ZeroCpy

Question 3 : Sort array


Concept

- Sort any array M of size d
- Using the two previous questions

Question 3 : Method



Question 3 : The execution time for different sizes of M



Batch merge

Batch merge : Concept

Concept

- A large number N of arrays A_i and B_i with $|A_i| + |B_i| = d \leq 1024$ for each $i = 1, \dots, N$.
- Merges two by two, each i , A_i and B_i .

Batch merge : structure

Structure used

- One big array M of size $N \cdot d$
- $M = [A_1|B_1|\dots|A_i|B_i|\dots|A_N|B_N]$
- $|A_i| + |B_i| = d \leq 1024$ for each $1 \leq i \leq N$
- $A_i = a_{i1}, a_{i2}, \dots, a_{il}$ and $B_i = b_{i1}, b_{i2}, \dots, b_{ik}$ with $l + k = d$
- Store sizes of A_i and B_i

Batch merge : index

```
int tidx = threadIdx.x % d;  
int Qt = (threadIdx.x - tidx) / d;  
int gbx = Qt + blockIdx.x * (blockDim.x / d);
```

Index


- tid_x : which diagonal does the thread takes care of.
- gb_x : index of the value of the size of the sub-array in the global array of sizes.
Allows us to retrieve the addresses of the sub-array A and B in M .

With gb_x , each block takes care of a sub-array.

With tid_x , each thread takes care of each diagonal.

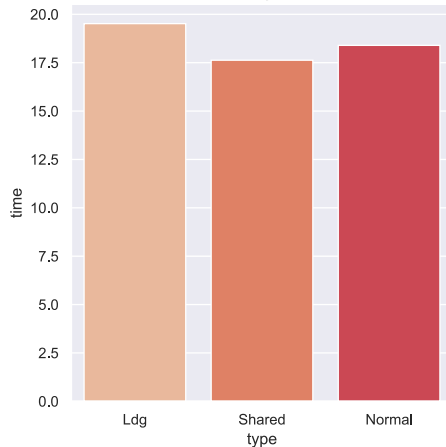
Batch merge : Different type of memory $d = 500$ and $N = 1000000$

Comparison of the execution time for mergeSmallBatch_k w.r.t the type of memory

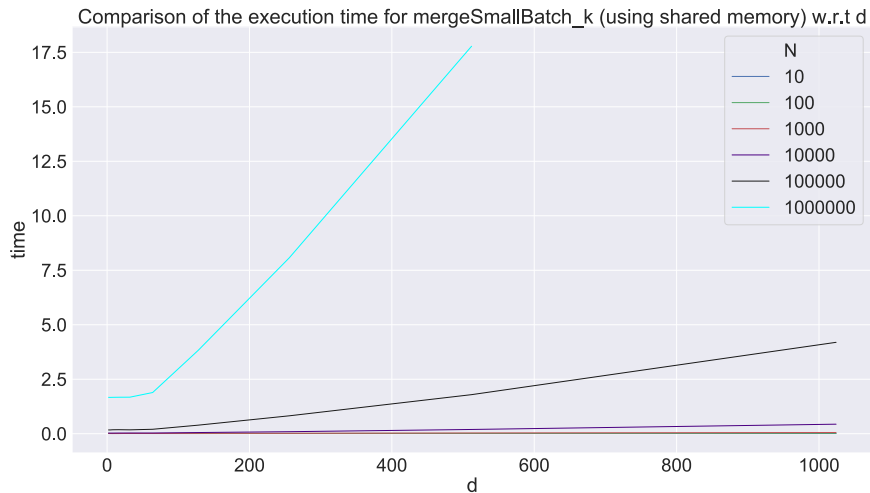


Batch merge : Different type of memory $d = 500$ and $N = 1000000$

Comparison of the execution time for mergeSmallBatch_k w.r.t the type of memory



Batch merge : The execution time for different d and N



A first Idea

Improving the parallel merge sort

Based on both questions, we came up with a new version of a fast merge sort.

It is based on the observation that in question 3, we have a nested loop which is very slow. Moreover, not every thread is used.

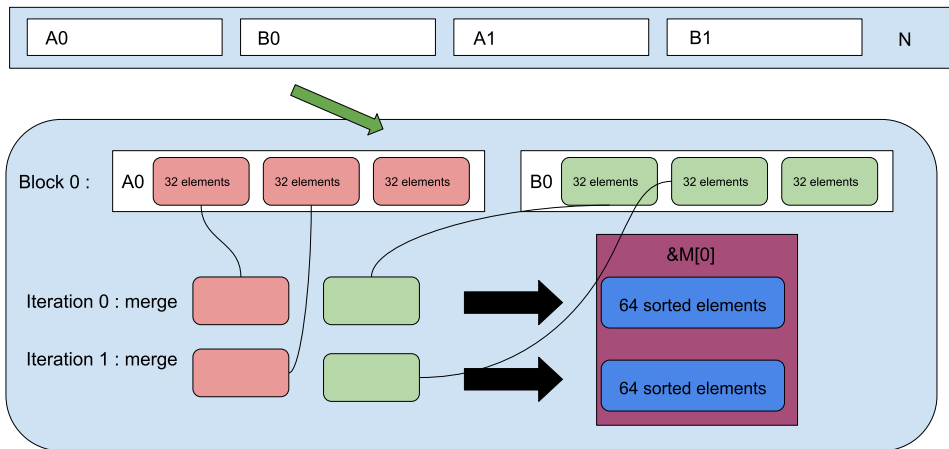
- 1 When $d \leq 1024$, `mergeSmallBatch` can be used in a single loop to get an array of sorted arrays of size 1024. Thus, the occupancy of the GPU is optimal since it is well schedule.
- 2 Then, a modified version of `pathBig_k` and `mergeSmall_k` will try to use the same strategy.

A first Idea

The strategy


- 1 Use block of size 32 or 64
- 2 Each block will merge start with a window of size `BlockDim`
- 3 `BlockDim` elements of A and B will be loaded sequentially in the shared memory
- 4 The path is computed and the merge is also done
- 5 Then, the action is repeated for every sub-arrays of size `BlockDim` in each arrays of size > 1024 .

Representation



Studying species of bacteria

Question 5 can be used to sort energy values of different "cells" that can be studied in a modelling system.



We can for example consider simple cells, interacting with each other. Each cell will have an **energy** counter, which will vary depending on the collisions and movements of the cell. We could want to know the cumulative distribution of the energy with regard to the cell kind.

Studying species of bacteria

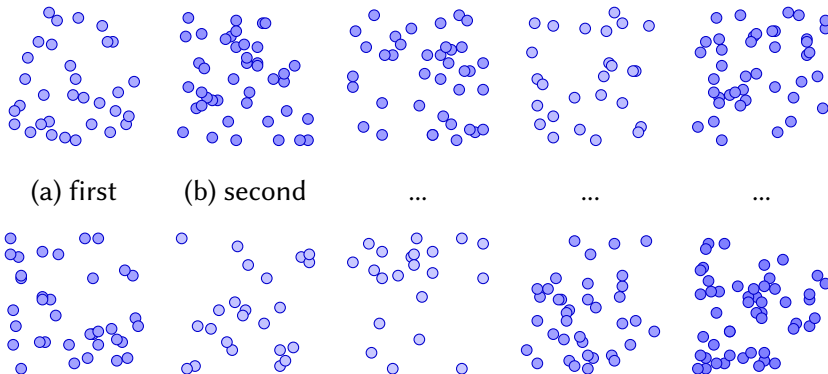
We thus need to compute the table containing the energy of each cell. We suppose that each cell's energy is automatically modified in the table. Then, we would end up with something close to :

Species/energy	0-50	50-100	> 100
A	10	5	0
B	60	1	1
⋮	⋮	⋮	⋮
C	30	30	30

Questions such as "Which species survived the best" or "What makes a species thrive?" can arise.

Launching different simulations in parallel

We could also want to study the same dynamics but on different populations set with different types of parameters. Again, if we study groups of size ≤ 1024 , we can compute every histogram at a fast pace.



Launching different simulations in parallel

`mergeSmallBatch` will be used on the following array

Speed/energy	0-50	50-100	> 100
4	10	5	0
6	60	1	1
⋮	⋮	⋮	⋮
10	30	30	30

We would loop over calls of `mergeSmallBatch`, changing d in order to merge and sort using the strategy from question 3. This is very efficient : at every step of the loop, each thread works on the same amount of values

Thanks for your attention

If you have any questions feel free to ask!