

Code source et prsentation du travail le 10/12/2020

La dure des presentations est de ~ 15 minutes par binme

1 Merge path and sort

[illegible]

1

Algorithm 1 Sequential Merge Path

Require: A and B are two sorted arrays

Ensure: M is the merged array of A and B with $|M| = |A| + |B|$

procedure MERGEPATH (A, B, M)

$j = 0$ and $i = 0$

while $i + j < |M|$ **do**

if $i \geq |A|$ **then**

$M[i+j] = B[j]$

$j = j + 1$

▷ The path goes right

else if $j \geq |B|$ or $A[i] < B[j]$ **then**

$M[i+j] = A[i]$

$i = i + 1$

▷ The path goes down

else

$M[i+j] = B[j]$

$j = j + 1$

▷ The path goes right

end if

end while

end procedure

Algorithm 2 Merge Path (Indexes of n threads are 0 to $n - 1$)

Require: A and B are two sorted arrays

Ensure: M is the merged array of A and B with $|M| = |A| + |B|$

for each thread **in parallel do**

$i = \text{index of the thread}$

if $i > |A|$ **then**

$K = (i - |A|, |A|)$

▷ Low point of diagonal

$P = (|A|, i - |A|)$

▷ High point of diagonal

else

$K = (0, i)$

$P = (i, 0)$

end if

while True **do**

$offset = \text{abs}(K_y - P_y)/2$

$Q = (K_x + offset, K_y - offset)$

if $Q_y \geq 0$ and $Q_x \leq B$ and

$(Q_y = |A|$ or $Q_x = 0$ or $A[Q_y] > B[Q_x - 1])$ **then**

if $Q_x = |B|$ or $Q_y = 0$ or $A[Q_y - 1] \leq B[Q_x]$ **then**

if $Q_y < |A|$ and $(Q_x = |B|$ or $A[Q_y] \leq B[Q_x])$ **then**

$M[i] = A[Q_y]$

▷ Merge in M

else

$M[i] = B[Q_x]$

end if

Break

else

$K = (Q_x + 1, Q_y - 1)$

end if

else

$P = (Q_x - 1, Q_y + 1)$

end if

end while

end for

Each point of the grid has a coordinate $(i, j) \in \llbracket 0, |A| \rrbracket \times \llbracket 0, |B| \rrbracket$. The merge path starts from the

point $(i, j) = (0, 0)$ on the left top corner of the grid. If $A[i] < B[j]$ the path goes down else it goes right. The array $\llbracket 0, |A| - 1 \rrbracket \times \llbracket 0, |B| - 1 \rrbracket$ of boolean values $A[i] < B[j]$ is not important in the algorithm. However, it shows clearly that the merge path is a frontier between ones and zeros.

To parallelize the algorithm, the grid has to be extended to the maximum size equal to $\max(|A|, |B|) \times \max(|A|, |B|)$. We denote K_0 and P_0 respectively the low point and the high point of the ascending diagonals Δ_k . On GPU, each thread $k \in \llbracket 0, |A| + |B| - 1 \rrbracket$ is responsible of one diagonal. It finds the intersection of the merge path and the diagonal Δ_k with a binary search described in Algorithm 2.

1. For $|A| + |B| \leq 1024$, write a kernel `mergeSmall.k` that merges A and B using only one block of threads.
2. For any size $|A| + |B| = d$ sufficiently smaller than the global memory, write two kernels that merge A and B using various blocks: The first kernel `pathBig.k` finds the merge path and the second one `mergeBig.k` merges A and B .
3. Looping on appropriate calls of `pathBig.k` and of `mergeBig.k`, write a function that sorts any array M of size d sufficiently smaller than the global memory. Give the execution time with respect to d .

2 Batch merge

In this part, we assume that we have a large number $N (\geq 1e3)$ of arrays $\{A_i\}_{1 \leq i \leq N}$ and $\{B_i\}_{1 \leq i \leq N}$ with $|A_i| + |B_i| = d \leq 1024$ for each i . Using some changes on `mergeSmall.k`, we would like to write `mergeSmallBatch.k` that merges two by two, for each i , A_i and B_i .

Given a fixed common size $d \leq 1024$, `mergeSmallBatch.k` is launched using the syntax

```
mergeSmallBatch.k<<<numBlocks, threadsPerBlock>>>(...);
```

with `threadsPerBlock` is multiple of d but smaller than 1024 and `numBlocks` is an arbitrary sufficiently big number.

4. Explain why the indices

```
int tid = threadIdx.x % d;
```

```
int Qt = (threadIdx.x - tid) / d;
```

```
int gbx = Qt + blockIdx.x * (blockDim.x / d);
```

are important in the definition of `mergeSmallBatch.k`.
5. Write the kernel `mergeSmallBatch.k` that batch merges two by two $\{A_i\}_{1 \leq i \leq N}$ and $\{B_i\}_{1 \leq i \leq N}$. Give the execution time with respect to $d = 4, 8, \dots, 1024$.

3 Merge sort applications

You are free to give a specific application of either question 3. or question 5.

References

- [1] O. Green, R. McColl and D. A. Bader GPU Merge Path - A GPU Merging Algorithm. *26th ACM International Conference on Supercomputing (ICS)*, San Servolo Island, Venice, Italy, June 25-29, 2012.